

AI Stack Development Lead

Job Description

Experience	5–8 years relevant experience
Contract Duration	[TBC]
Location	UAE (on-site)
Platform Area	AI Platforms – Azure

About the Role

The AI Stack Development Lead will design, build, and lead the development of the enterprise AI application stack on Microsoft Azure, turning approved AI use cases into reliable, secure, production-grade AI services and agents. The role spans solution architecture, hands-on development, productionisation, and technical leadership of a small AI engineering team.

This is a hands-on lead role. It owns the engineering standards, reusable patterns, and delivery of AI solutions built on Azure OpenAI, Azure AI Foundry, Azure AI Services, and MCP-based agent integrations, working closely with AI Platforms, Cloud Platforms, Data Platforms, Architecture, Security, and business teams.

It ensures AI solutions and agents are well-engineered, governed, observable, cost-efficient, and ready for enterprise adoption, covering corporate AI use cases, centralised productionisation, and function-specific agent building.

Key Responsibilities

- Lead the design and development of AI services, AI agents, retrieval-augmented generation (RAG) solutions, and copilots on Azure OpenAI, Azure AI Foundry, and Azure AI Services.
- Own the AI application stack architecture, including orchestration frameworks, MCP gateways and tool integrations, APIs, connectors, and the retrieval layer on Azure AI Search.
- Set and enforce engineering standards, reusable patterns, code quality, and design reviews across AI developers.
- Lead and mentor a small team of AI developers; plan sprints, estimate work, and manage technical risk and delivery quality.
- Build and support AI agents that connect securely to approved enterprise data sources, tools, APIs, and systems.
- Drive centralised productionisation of AI use cases, including evaluation, testing, release readiness, monitoring, documentation, and operational handover.
- Establish LLMOps practices with the AI Platforms team: prompt and model versioning, automated evaluation, regression testing, CI/CD, and infrastructure as code.
- Implement responsible AI, security, and governance controls by design, including content safety, prompt-injection defences, PII handling, identity and access, and auditability.
- Work with business teams to assess, prioritise, and onboard AI use cases, and translate them into technical designs and delivery plans.
- Evaluate models, frameworks, and tools; run proofs of concept and provide roadmap inputs.
- Monitor solution quality, usage, cost, and performance, and resolve application, integration, and data-access issues.
- Maintain technical documentation, reusable implementation patterns, and enablement material for AI developers.

Required Qualifications & Experience

- Bachelor's degree in Computer Science, Engineering, AI, Data Science, or a related discipline.

- 5–8 years of hands-on software or AI engineering experience, including at least 2 years building production Generative AI or agent-based solutions on Microsoft Azure.
- Proven experience leading a small development team or acting as technical lead on AI/ML delivery.
- Strong Python; working knowledge of C#/ .NET or TypeScript.
- Hands-on experience with Azure OpenAI, Azure AI Foundry (models, Agent Service, evaluations), Azure AI Search, and Azure AI Services.
- Experience with agent and orchestration frameworks (Semantic Kernel, Microsoft Agent Framework, LangChain/LangGraph, AutoGen) and MCP-based tool integration.
- Experience building and securing APIs and integrations on Azure (API Management, Functions, Container Apps or AKS, Cosmos DB, Key Vault, Entra ID).
- Experience with CI/CD, infrastructure as code (Bicep or Terraform), and productionising AI solutions with evaluation, monitoring, and cost controls.
- Working understanding of responsible AI, data governance, identity, permissions, and security controls for AI systems.
- Strong analytical, problem-solving, and communication skills; able to work with technical teams and business stakeholders.

Preferred Skills

- Experience with AI coding assistants such as Claude Code, OpenAI Codex, GitHub Copilot, and Cursor, and enabling their adoption across engineering teams.
- Experience with Microsoft 365 Copilot extensibility, Copilot Studio, and Power Platform.
- Exposure to Azure Machine Learning, Azure Databricks, or Microsoft Fabric, and to enterprise data platforms such as Snowflake, Azure Data Factory, or Power BI.
- Experience fine-tuning or hosting open-weight models (Azure AI Foundry model catalogue, GPU compute, vLLM or ONNX Runtime).
- Familiarity with observability and evaluation for LLM applications (OpenTelemetry, Application Insights, Azure AI Foundry evaluations, RAGAS or DeepEval).
- Knowledge of AI governance frameworks and regulation (NIST AI RMF, ISO/IEC 42001, EU AI Act) and the UAE Personal Data Protection Law (PDPL).
- Experience with front-end frameworks (React or Next.js) or rapid prototyping tools (Streamlit, Gradio) for AI experiences.
- Prior experience delivering enterprise AI solutions in the UAE or GCC region.

Preferred Certifications

- Microsoft Certified: Azure AI Engineer Associate (AI-102).
- Microsoft Certified: Azure Solutions Architect Expert (AZ-305) or Azure Developer Associate (AZ-204).
- Microsoft Certified: DevOps Engineer Expert (AZ-400) or Azure Data Scientist Associate (DP-100).
- AI, cloud, or security-related certifications are an advantage.

Key Success Measures

- A well-architected, reusable AI application stack on Azure adopted across AI use cases.
- AI use cases delivered from prototype to production with consistent quality, evaluation, and monitoring.
- Reliable, secure, and governed operation of AI services and agents.
- High engineering standards, code quality, and reuse across the AI developer team.
- Timely resolution of application, integration, and data-access issues.
- Demonstrated compliance with responsible AI, security, and operational standards.